Take a `moderndive` into introductory linear regression with R

Albert Y. Kim[1], Chester Ismay[2], and Max Kuhn[3]

1 Assistant Professor of Statistical and Data Sciences, Smith College, Northampton, MA, USA. **2** Data Science Evangelist, DataRobot, Portland, OR, USA. **3** Software Engineer, RStudio, USA.

Introduction

Linear regression has long been a staple of introductory statistics courses. While the curricula of introductory statistics courses has much evolved of late, the overall importance of regression remains the same (American Statistical Association Undergraduate Guidelines Workgroup 2016). To facilitate the teaching of regression while leveraging modern computational tools, we've created the `moderndive` R package of datasets and functions for tidyverse-friendly introductory linear regression (Wickham, Averick, et al. 2019).

This R package is designed to supplement the book "Statistical Inference via Data Science: A ModernDive into R and the Tidyverse" (Ismay and Kim 2019). Note that the book is also available online at https://www.moderndive.com and is referred to as "ModernDive" for short.

Before we proceed, let's load all the R packages we are going to need.

```r
library(moderndive)
library(ggplot2)
library(dplyr)
library(readr)
library(knitr)
library(broom)
```

Let's consider data gathered from end of semester student evaluations for a sample of 463 courses taught by 94 professors from the University of Texas at Austin (Diez, Barr, and Çetinkaya-Rundel 2015). This data is included in the `evals` data frame from the `moderndive` package.

In the following table, we present a subset of 9 of the 14 variables included for a random sample of 5 courses[1]:

1. `ID` uniquely identifies the course whereas `prof_ID` identifies the professor who taught this course. This distinction is important since many professors taught more than one course.
2. `score` is the outcome variable of interest: average professor evaluation score out of 5 as given by the students in this course.
3. The remaining variables are demographic variables describing that course's instructor, including `bty_avg` (average "beauty" score) for that professor as given by a panel of 6 students.[2]

---

[1]For details on the remaining 5 variables, see the help file by running `?evals`.
[2]Note that `gender` was collected as a binary variable at the time of the study (2005).

ID	prof_ID	score	age	bty_avg	gender	ethnicity	language	rank
129	23	3.7	62	3.000	male	not minority	english	tenured
109	19	4.7	46	4.333	female	not minority	english	tenured
28	6	4.8	62	5.500	male	not minority	english	tenured
434	88	2.8	62	2.000	male	not minority	english	tenured
330	66	4.0	64	2.333	male	not minority	english	tenured

Regression analysis the "good old-fashioned" way

Let's fit a simple linear regression model of teaching `score` as a function of instructor `age` using the `lm()` function.

```
score_model <- lm(score ~ age, data = evals)
```

Let's now study the output of the fitted model `score_model` "the good old-fashioned way": using `summary()` which calls `summary.lm()` behind the scenes (we'll refer to them interchangeably throughout this paper).

```
summary(score_model)
##
## Call:
## lm(formula = score ~ age, data = evals)
##
## Residuals:
##     Min      1Q  Median      3Q     Max
## -1.9185 -0.3531  0.1172  0.4172  0.8825
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  4.461932   0.126778  35.195   <2e-16 ***
## age         -0.005938   0.002569  -2.311   0.0213 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5413 on 461 degrees of freedom
## Multiple R-squared:  0.01146,    Adjusted R-squared:  0.009311
## F-statistic: 5.342 on 1 and 461 DF,  p-value: 0.02125
```

Here are 5 common student questions we've heard over the years in our introductory statistics courses based on this output:

1. "Wow! Look at those p-value stars! Stars are good, so I should try to get many stars, right?"
2. "How do we extract the values in the regression table?"
3. "Where are the fitted/predicted values and residuals?"
4. "How do I apply this model to a new set of data to make predictions?"
5. "What is all this other stuff at the bottom?"

Regression analysis using `moderndive`

To address these questions, we've included three functions in the `moderndive` package that take a fitted model object as input and return the same information as `summary.lm()`, but output them in tidyverse-friendly format (Wickham, Averick, et al. 2019). As we'll see in subsection 3.1, while these three functions are merely wrappers to existing functions in

the `broom` package for converting statistical objects into tidy tibbles, we modified them with the introductory statistics student in mind (Robinson and Hayes 2019).

1. Get a tidy regression table **with confidence intervals**:

```
get_regression_table(score_model)
## # A tibble: 2 x 7
##   term      estimate std_error statistic p_value lower_ci upper_ci
##   <chr>        <dbl>     <dbl>     <dbl>   <dbl>    <dbl>    <dbl>
## 1 intercept    4.46     0.127     35.2     0        4.21     4.71
## 2 age         -0.006    0.003     -2.31    0.021   -0.011   -0.001
```

2. Get information on each point/observation in your regression, including fitted/predicted values and residuals, in a single data frame:

```
get_regression_points(score_model)
## # A tibble: 463 x 5
##       ID score   age score_hat residual
##    <int> <dbl> <int>     <dbl>    <dbl>
## 1      1   4.7    36      4.25    0.452
## 2      2   4.1    36      4.25   -0.148
## 3      3   3.9    36      4.25   -0.348
## 4      4   4.8    36      4.25    0.552
## 5      5   4.6    59      4.11    0.488
## 6      6   4.3    59      4.11    0.188
## 7      7   2.8    59      4.11   -1.31
## 8      8   4.1    51      4.16   -0.059
## 9      9   3.4    51      4.16   -0.759
## 10    10   4.5    40      4.22    0.276
## # ... with 453 more rows
```

3. Get scalar summaries of a regression fit including R^2 and R^2_{adj} but also the (root) mean-squared error:

```
get_regression_summaries(score_model)
## # A tibble: 1 x 8
##   r_squared adj_r_squared   mse  rmse sigma statistic p_value    df
##       <dbl>         <dbl> <dbl> <dbl> <dbl>     <dbl>   <dbl> <dbl>
## 1     0.011         0.009 0.292 0.540 0.541      5.34   0.021     2
```

Bonus: Visualizing parallel slopes models with `moderndive`

Furthermore, say you would like to create a visualization of the relationship between two numerical variables and a third categorical variable with k levels. Let's create this using a colored scatterplot via the `ggplot2` package for data visualization (Wickham, Chang, et al. 2019). Using `geom_smooth(method = "lm", se = FALSE)` yields a visualization of an *interaction model* where each of the k regression lines has their own intercept and slope. For example in Figure 1, we extend our previous regression model by now mapping the categorical variable `ethnicity` to the `color` aesthetic.

```
# Code to visualize interaction model:
ggplot(evals, aes(x = age, y = score, color = ethnicity)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE) +
  labs(x = "Age", y = "Teaching score", color = "Ethnicity")
```

However, many introductory statistics courses start with the easier to teach "common

Figure 1: Visualization of interaction model.

slope, different intercepts" regression model, also known as the *parallel slopes* model. However, no argument to plot such models exists within `geom_smooth()`.

Evgeni Chasnovski thus wrote a custom `geom_` extension to `ggplot2` called `geom_parallel_slopes()`; this extension is included in the `moderndive` package. Much like `geom_smooth()` from the `ggplot2` package, you merely add a `geom_parallel_slopes()` layer to the code, resulting in Figure 2.

```r
# Code to visualize parallel slopes model:
ggplot(evals, aes(x = age, y = score, color = ethnicity)) +
  geom_point() +
  geom_parallel_slopes(se = FALSE) +
  labs(x = "Age", y = "Teaching score", color = "Ethnicity")
```

At this point however, students will inevitably ask a sixth question: "When would you ever use a parallel slopes model?"

Why should you use the `moderndive` package?

To recap this introduction, we believe that the following functions included in the `moderndive` package

1. `get_regression_table()`
2. `get_regression_points()`
3. `get_regression_summaries()`
4. `geom_parallel_slopes()`

are effective pedagogical tools that can help address the six common questions posed by students about introductory linear regression performed in R. We now argue why.

Figure 2: Visualization of parallel slopes model.

Features

1. Less p-value stars, more confidence intervals

The first common student question:

> "Wow! Look at those p-value stars! Stars are good, so I should try to get many stars, right?"

We argue that the `summary.lm()` output is deficient in an introductory statistics setting because:

1. The `Signif. codes: 0 '' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1` only encourage **p-hacking**. In case you have not yet been convinced of the perniciousness of p-hacking, perhaps comedian John Oliver can convince you.

2. While not a silver bullet for eliminating misinterpretations of statistical inference, confidence intervals present students with a sense of the associated effect sizes of any explanatory variables. Thus, practical as well as statistical significance is emphasized. These are not included by default in the output of `summary.lm()`.

Instead of `summary()`, let's use the `get_regression_table()` function:

```
get_regression_table(score_model)
## # A tibble: 2 x 7
##   term       estimate std_error statistic p_value lower_ci upper_ci
##   <chr>         <dbl>     <dbl>     <dbl>   <dbl>    <dbl>    <dbl>
## 1 intercept      4.46     0.127      35.2   0         4.21     4.71
## 2 age           -0.006    0.003      -2.31   0.021    -0.011   -0.001
```

Observe how the p-value stars are omitted and confidence intervals for the point estimates of all regression parameters are included by default. By including them in the output, we can easily emphasize to students that they "surround" the point estimates in the `estimate` column. Note the confidence level is defaulted to 95%.

2. Outputs as tibbles

The second common student question:

> "How do we extract the values in the regression table?"

While one might argue that extracting the intercept and slope coefficients can be simply done using `coefficients(score_model)`, what about the standard errors? For example, a Google query of "*how do I extract standard errors from lm in R*" yielded results from the R mailing list and from crossvalidated suggesting we run:

```
sqrt(diag(vcov(score_model)))
## (Intercept)         age
## 0.126778499 0.002569157
```

We argue that this task shouldn't be this hard, especially in an introductory statistics setting. To rectify this, the three `get_regression_X()` functions all return data frames in the tidyverse-style tibble (tidy table) format (Müller and Wickham 2019). Therefore you can easily extract columns using the `pull()` function from the `dplyr` package (Wickham et al. 2020):

```
get_regression_table(score_model) %>%
  pull(std_error)
## [1] 0.127 0.003
```

or equivalently you can use the `$` sign operator from base R:

```
get_regression_table(score_model)$std_error
## [1] 0.127 0.003
```

Furthermore, by piping the above `get_regression_table(score_model)` output into the `kable()` function from the `knitr` package (Xie 2020), you can obtain aesthetically pleasing regression tables in R Markdown documents, instead of tables written in jarring computer output font:

```
get_regression_table(score_model) %>%
  kable()
```

term	estimate	std_error	statistic	p_value	lower_ci	upper_ci
intercept	4.462	0.127	35.195	0.000	4.213	4.711
age	-0.006	0.003	-2.311	0.021	-0.011	-0.001

3. Birds of a feather should flock together: Fitted values & residuals

The third common student question:

> "Where are the fitted/predicted values and residuals?"

How can we extract point-by-point information from a regression model, such as the fitted/predicted values and the residuals? (Note we only display the first 10 out of 463 of such values for brevity's sake.)

```
fitted(score_model)

##        1        2        3        4        5        6        7
## 4.248156 4.248156 4.248156 4.248156 4.111577 4.111577 4.111577
##        8        9       10
## 4.159083 4.159083 4.224403
```

```
residuals(score_model)
```

```
##          1           2           3           4           5
##  0.45184376 -0.14815624 -0.34815624  0.55184376  0.48842294
##          6           7           8           9          10
##  0.18842294 -1.31157706 -0.05908286 -0.75908286  0.27559666
```

But why have the original explanatory/predictor `age` and outcome variable `score` in `evals`, the fitted/predicted values `score_hat`, and `residual` floating around in separate vectors? Since each observation relates to the same course, we argue it makes more sense to organize them together in the same data frame using `get_regression_points()`:

```
score_model_points <- get_regression_points(score_model)
score_model_points
```

```
## # A tibble: 10 x 5
##       ID score   age score_hat residual
##    <int> <dbl> <int>     <dbl>    <dbl>
##  1     1   4.7    36      4.25    0.452
##  2     2   4.1    36      4.25   -0.148
##  3     3   3.9    36      4.25   -0.348
##  4     4   4.8    36      4.25    0.552
##  5     5   4.6    59      4.11    0.488
##  6     6   4.3    59      4.11    0.188
##  7     7   2.8    59      4.11   -1.31
##  8     8   4.1    51      4.16   -0.059
##  9     9   3.4    51      4.16   -0.759
## 10    10   4.5    40      4.22    0.276
```

Observe that the original outcome variable `score` and explanatory/predictor variable `age` are now supplemented with the fitted/predicted values `score_hat` and `residual` columns. By putting the fitted values, predicted values, and residuals next to the original data, we argue that the computation of these values is less opaque. For example, instructors can easily emphasize how all values in the first row of output are computed.

Furthermore, recall that since all outputs in the `moderndive` package are tibble data frames, custom residual analysis plots can be created instead of relying on the default plots yielded by `plot.lm()`. For example, we can check for the normality of residuals using the histogram of residuals shown in Figure 3.

```
# Code to visualize distribution of residuals:
ggplot(score_model_points, aes(x = residual)) +
  geom_histogram(bins = 20) +
  labs(x = "Residual", y = "Count")
```

As another example, we can investigate potential relationships between the residuals and all explanatory/predictor variables and the presence of heteroskedasticity using partial residual plots, like the partial residual plot over age shown in Figure 4. If the term "heteroskedasticity" is new to you, it corresponds to the variability of one variable being unequal across the range of values of another variable. The presence of heteroskedasticity violates one of the assumptions of inference for linear regression.

```
# Code to visualize partial residual plot over age:
ggplot(score_model_points, aes(x = age, y = residual)) +
  geom_point() +
  labs(x = "Age", y = "Residual")
```
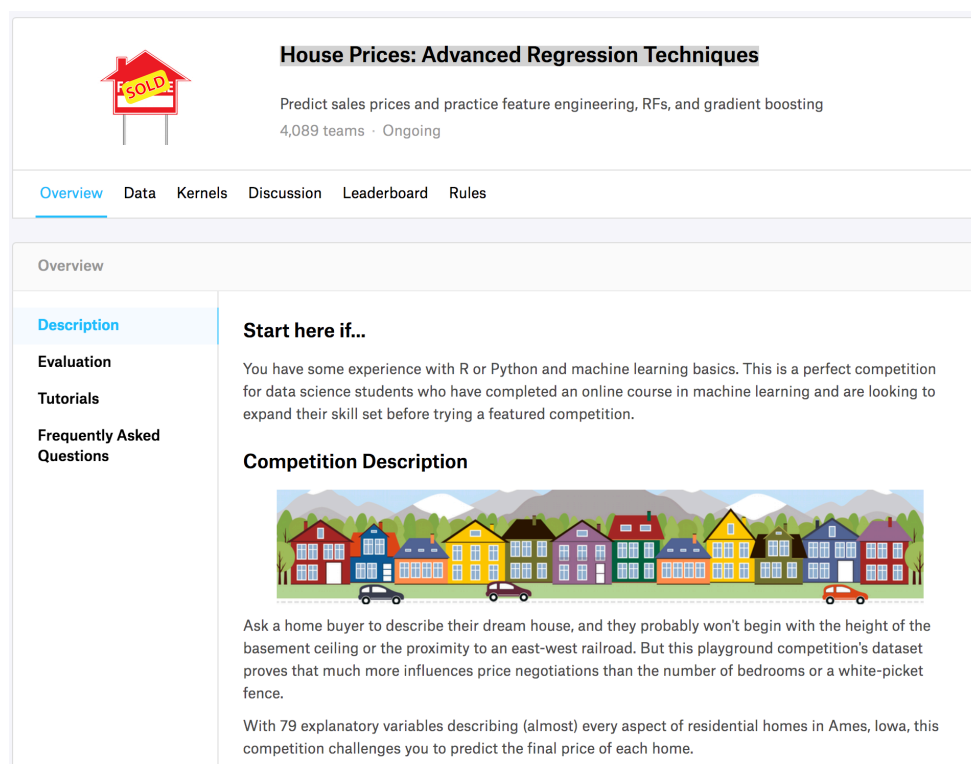
Figure 3: Histogram visualizing distribution of residuals.

Figure 4: Partial residual residual plot over age.

, (). Take a `moderndive` into introductory linear regression with R. *Journal of Open Source Software*, (), . https://doi.org/

Figure 5: House prices Kaggle competition homepage.

4. A quick-and-easy Kaggle predictive modeling competition submission!

The fourth common student question:

> "How do I apply this model to a new set of data to make predictions?"

With the fields of machine learning and artificial intelligence gaining prominence, the importance of predictive modeling cannot be understated. Therefore, we've designed the `get_regression_points()` function to allow for a `newdata` argument to quickly apply a previously fitted model to new observations.

Let's create an artificial "new" dataset consisting of two instructors of age 39 and 42 and save it in a tibble data frame called `new_prof`. We then set the `newdata` argument to `get_regression_points()` to apply our previously fitted model `score_model` to this new data, where `score_hat` holds the corresponding fitted/predicted values.

```
new_prof <- tibble(age = c(39, 42))
get_regression_points(score_model, newdata = new_prof)
## # A tibble: 2 x 3
##      ID   age score_hat
##   <int> <dbl>     <dbl>
## 1     1    39      4.23
## 2     2    42      4.21
```

Let's do another example, this time using the Kaggle House Prices: Advanced Regression Techniques practice competition (Figure 5 displays the homepage for this competition).

This Kaggle competition requires you to fit/train a model to the provided `train.csv` training set to make predictions of house prices in the provided `test.csv` test set. We

Figure 6: Resulting Kaggle RMSLE score.

present an application of the `get_regression_points()` function allowing students to participate in this Kaggle competition. It will:

1. Read in the training and test data.
2. Fit a naive model of house sale price as a function of year sold to the training data.
3. Make predictions on the test data and write them to a `submission.csv` file that can be submitted to Kaggle using `get_regression_points()`. Note the use of the `ID` argument to use the `id` variable in `test` to identify the rows (a requirement of Kaggle competition submissions).

```r
library(readr)
library(dplyr)
library(moderndive)

# Load in training and test set
train <- read_csv("https://moderndive.com/data/train.csv")
test <- read_csv("https://moderndive.com/data/test.csv")

# Fit model:
house_model <- lm(SalePrice ~ YrSold, data = train)

# Make predictions and save in appropriate data frame format:
submission <- house_model %>%
  get_regression_points(newdata = test, ID = "Id") %>%
  select(Id, SalePrice = SalePrice_hat)

# Write predictions to csv:
write_csv(submission, "submission.csv")
```

After submitting `submission.csv` to the leaderboard for this Kaggle competition, we obtain a "root mean squared logarithmic error" (RMSLE) score of 0.42918 as seen in Figure 6.

5. Scalar summaries of linear regression model fits

The fifth common student question:

> "What is all this other stuff at the bottom?"

Recall the output of the standard `summary.lm()` from earlier:

```
## 
## Call:
## lm(formula = score ~ age, data = evals)
## 
## Residuals:
##     Min      1Q  Median      3Q     Max
## -1.9185 -0.3531  0.1172  0.4172  0.8825
## 
```

Figure 7: Interaction (left) and parallel slopes (right) models.

```
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  4.461932   0.126778  35.195   <2e-16 ***
## age         -0.005938   0.002569  -2.311   0.0213 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5413 on 461 degrees of freedom
## Multiple R-squared:  0.01146,    Adjusted R-squared:  0.009311
## F-statistic: 5.342 on 1 and 461 DF,  p-value: 0.02125
```

Say we wanted to extract the scalar model summaries at the bottom of this output, such as R^2, R^2_{adj}, and the F-statistic. We can do so using the `get_regression_summaries()` function.

```
get_regression_summaries(score_model)
## # A tibble: 1 x 8
##   r_squared adj_r_squared   mse  rmse sigma statistic p_value    df
##       <dbl>         <dbl> <dbl> <dbl> <dbl>     <dbl>   <dbl> <dbl>
## 1     0.011         0.009 0.292 0.540 0.541      5.34   0.021     2
```

We've supplemented the standard scalar summaries output yielded by `summary()` with the mean squared error `mse` and root mean squared error `rmse` given their popularity in machine learning settings.

6. Plot parallel slopes regression models

Finally, the last common student question:

> "When would you ever use a parallel slopes model?"

For example, recall the earlier visualizations of the interaction and parallel slopes models for teaching score as a function of age and ethnicity we saw in Figures 1 and 2. Let's present both visualizations side-by-side in Figure 7.

Students might be wondering "Why would you use the parallel slopes model on the right when the data clearly form an"X" pattern as seen in the interaction model on the left?" This is an excellent opportunity to gently introduce the notion of *model selection* and *Occam's Razor*: an interaction model should only be used over a parallel slopes model **if**

the additional complexity of the interaction model is warranted. Here, we define model "complexity/simplicity" in terms of the number of parameters in the corresponding regression tables:

```
# Regression table for interaction model:
interaction_evals <- lm(score ~ age * ethnicity, data = evals)
get_regression_table(interaction_evals)
## # A tibble: 4 x 7
##   term        estimate std_error statistic p_value lower_ci upper_ci
##   <chr>          <dbl>     <dbl>     <dbl>   <dbl>    <dbl>    <dbl>
## 1 intercept       2.61     0.518      5.04   0         1.59     3.63
## 2 age             0.032    0.011      2.84   0.005     0.01     0.054
## 3 ethnicitynot~   2.00     0.534      3.74   0         0.945    3.04
## 4 age:ethnicit~  -0.04     0.012     -3.51   0        -0.063   -0.018


# Regression table for parallel slopes model:
parallel_slopes_evals <- lm(score ~ age + ethnicity, data = evals)
get_regression_table(parallel_slopes_evals)
## # A tibble: 3 x 7
##   term        estimate std_error statistic p_value lower_ci upper_ci
##   <chr>          <dbl>     <dbl>     <dbl>   <dbl>    <dbl>    <dbl>
## 1 intercept       4.37     0.136     32.1    0         4.1      4.63
## 2 age            -0.006    0.003     -2.5    0.013    -0.012   -0.001
## 3 ethnicitynot~   0.138    0.073      1.89   0.059    -0.005    0.282
```

The interaction model is "more complex" as evidenced by its regression table involving 4 rows of parameter estimates whereas the parallel slopes model is "simpler" as evidenced by its regression table involving only 3 parameter estimates. It can be argued however that this additional complexity is warranted given the clearly different slopes in the left-hand plot of Figure 7.

We now present a contrasting example, this time from Chapter 6 of the online version of ModernDive Subsection 6.3.1 involving Massachusetts USA public high schools.[3] Let's plot both the interaction and parallel slopes models in Figure 8.

```
# Code to plot interaction and parallel slopes models for MA_schools
ggplot(MA_schools,
       aes(x = perc_disadvan, y = average_sat_math, color = size)) +
  geom_point(alpha = 0.25) +
  labs(x = "% economically disadvantaged",
       y = "Math SAT Score",
       color = "School size") +
  geom_smooth(method = "lm", se = FALSE)

ggplot(MA_schools,
       aes(x = perc_disadvan, y = average_sat_math, color = size)) +
  geom_point(alpha = 0.25) +
  labs(x = "% economically disadvantaged",
       y = "Math SAT Score",
       color = "School size") +
  geom_parallel_slopes(se = FALSE)
```

In terms of the corresponding regression tables, observe that the corresponding regression table for the parallel slopes model has 4 rows as opposed to the 6 for the interaction model, reflecting its higher degree of "model simplicity."

---

[3]For more details on this dataset, see the help file by running `?MA_schools`.

Figure 8: Interaction (left) and parallel slopes (right) models.

```
# Regression table for interaction model:
interaction_MA <-
  lm(average_sat_math ~ perc_disadvan * size, data = MA_schools)
get_regression_table(interaction_MA)
## # A tibble: 6 x 7
##   term          estimate std_error statistic p_value lower_ci upper_ci
##   <chr>            <dbl>     <dbl>     <dbl>   <dbl>    <dbl>    <dbl>
## 1 intercept        594.      13.3      44.7       0     568.     620.
## 2 perc_disadvan   -2.93     0.294     -9.96       0    -3.51    -2.35
## 3 sizemedium      -17.8      15.8     -1.12   0.263    -48.9     13.4
## 4 sizelarge       -13.3      13.8    -0.962   0.337    -40.5     13.9
## 5 perc_disadva~   0.146     0.371     0.393   0.694   -0.585    0.877
## 6 perc_disadva~   0.189     0.323     0.586   0.559   -0.446    0.824

# Regression table for parallel slopes model:
parallel_slopes_MA <-
  lm(average_sat_math ~ perc_disadvan + size, data = MA_schools)
get_regression_table(parallel_slopes_MA)
## # A tibble: 4 x 7
##   term          estimate std_error statistic p_value lower_ci upper_ci
##   <chr>            <dbl>     <dbl>     <dbl>   <dbl>    <dbl>    <dbl>
## 1 intercept        588.      7.61      77.3       0     573.     603.
## 2 perc_disadvan   -2.78     0.106     -26.1       0    -2.99    -2.57
## 3 sizemedium      -11.9      7.54     -1.58   0.115    -26.7     2.91
## 4 sizelarge       -6.36      6.92    -0.919   0.359    -20.0     7.26
```

Unlike our earlier comparison of interaction and parallel slopes models in Figure 7, in this case it could be argued that the additional complexity of the interaction model is *not* warranted since the 3 three regression lines in the left-hand interaction are already somewhat parallel. Therefore the simpler parallel slopes model should be favored.

Going one step further, notice how the three regression lines in the visualization of the parallel slopes model in the right-hand plot of Figure 8 have similar intercepts. In can thus be argued that the additional model complexity induced by introducing the categorical variable school `size` is not warranted. Therefore, a simple linear regression model using only `perc_disadvan` percent of the student body that are economically disadvantaged should be favored.

While many students will inevitably find these results depressing, in our opinion, it is important to additionally emphasize that such regression analyses can be used as an empowering tool to bring to light inequities in access to education and inform policy decisions.

The Details

Three wrappers to `broom` functions

As we mentioned earlier, the three `get_regression_*` functions are merely wrappers of functions from the `broom` package for converting statistical analysis objects into tidy tibbles along with a few added tweaks, but with the introductory statistics student in mind (Robinson and Hayes 2019):

1. `get_regression_table()` is a wrapper for `broom::tidy()`.
2. `get_regression_points()` is a wrapper for `broom::augment()`.
3. `get_regression_summaries` is a wrapper for `broom::glance()`.

Why did we take this approach to address the initial 5 common student questions at the outset of the article?

1. By writing wrappers to pre-existing functions, instead of creating new custom functions, there is minimal re-inventing of the wheel necessary.
2. In our experience, novice R users had a hard time understanding the `broom` package function names `tidy()`, `augment()`, and `glance()`. To make them more user-friendly, the `moderndive` package wrappers have much more intuitively named `get_regression_table()`, `get_regression_points()`, and `get_regression_summaries()`.
3. The variables included in the outputs of the above 3 `broom` functions are not all applicable to an introductory statistics students and of those that were, we found them to be unintuitively named. We therefore cut out some of the variables from the output and renamed some of the remaining variables. For example, compare the outputs of the `get_regression_points()` wrapper function and the parent `broom::augment()` function.

```
get_regression_points(score_model)
## # A tibble: 463 x 5
##       ID score   age score_hat residual
##    <int> <dbl> <int>     <dbl>    <dbl>
## 1      1   4.7    36      4.25    0.452
## 2      2   4.1    36      4.25   -0.148
## 3      3   3.9    36      4.25   -0.348
## 4      4   4.8    36      4.25    0.552
## 5      5   4.6    59      4.11    0.488
## 6      6   4.3    59      4.11    0.188
## 7      7   2.8    59      4.11   -1.31
## 8      8   4.1    51      4.16   -0.059
## 9      9   3.4    51      4.16   -0.759
## 10    10   4.5    40      4.22    0.276
## # ... with 453 more rows
broom::augment(score_model)
## # A tibble: 463 x 9
##    score   age .fitted .se.fit  .resid    .hat .sigma  .cooksd
##    <dbl> <int>   <dbl>   <dbl>   <dbl>   <dbl>  <dbl>    <dbl>
## 1    4.7    36    4.25  0.0405   0.452 0.00560  0.542 1.97e-3
```

, (). Take a moderndive into introductory linear regression with R. *Journal of Open Source Software*, (), . https://doi.org/   14

```
##  2   4.1    36    4.25  0.0405 -0.148  0.00560  0.542 2.12e-4
##  3   3.9    36    4.25  0.0405 -0.348  0.00560  0.542 1.17e-3
##  4   4.8    36    4.25  0.0405  0.552  0.00560  0.541 2.94e-3
##  5   4.6    59    4.11  0.0371  0.488  0.00471  0.541 1.93e-3
##  6   4.3    59    4.11  0.0371  0.188  0.00471  0.542 2.88e-4
##  7   2.8    59    4.11  0.0371 -1.31   0.00471  0.538 1.39e-2
##  8   4.1    51    4.16  0.0261 -0.0591 0.00232  0.542 1.39e-5
##  9   3.4    51    4.16  0.0261 -0.759  0.00232  0.541 2.29e-3
## 10   4.5    40    4.22  0.0331  0.276  0.00374  0.542 4.88e-4
## # ... with 453 more rows, and 1 more variable: .std.resid <dbl>
```

The source code for these three `get_regression_*` functions can be found on GitHub.

Custom geometries

The `geom_parallel_slopes()` is a custom built `geom` extension to the `ggplot2` package. For example, the `ggplot2` webpage page gives instructions on how to create such extensions. The source code for `geom_parallel_slopes()` written by Evgeni Chasnovski can be found on GitHub.

Acknowledgements

References

American Statistical Association Undergraduate Guidelines Workgroup. 2016. "'Guidelines for Assessment and Instruction in Statistics Education (GAISE) in Statistics Education (GAISE) College Report College Report 2016'." Alexandria, VA: American Statistical Association. http://www.amstat.org/education/gaise.

Diez, D.M., C.D. Barr, and M. Çetinkaya-Rundel. 2015. *OpenIntro Statistics*. OpenIntro, Incorporated. https://books.google.com/books?id=xNMWswEACAAJ.

Ismay, Chester., and Albert Y. Kim. 2019. *Statistical Inference via Data Science: A Moderndive into R and the Tidyverse*. Chapman & Hall/Crc the R Series. CRC Press. https://moderndive.com/.

Müller, Kirill, and Hadley Wickham. 2019. *Tibble: Simple Data Frames*. https://CRAN.R-project.org/package=tibble.

Robinson, David, and Alex Hayes. 2019. *Broom: Convert Statistical Analysis Objects into Tidy Tibbles*. https://CRAN.R-project.org/package=broom.

Wickham, Hadley, Mara Averick, Jennifer Bryan, Winston Chang, Lucy D'Agostino McGowan, Romain François, Garrett Grolemund, et al. 2019. "Welcome to the tidyverse." *Journal of Open Source Software* 4 (43): 1686. https://doi.org/10.21105/joss.01686.

Wickham, Hadley, Winston Chang, Lionel Henry, Thomas Lin Pedersen, Kohske Takahashi, Claus Wilke, Kara Woo, and Hiroaki Yutani. 2019. *Ggplot2: Create Elegant Data*

Visualisations Using the Grammar of Graphics. https://CRAN.R-project.org/package=ggplot2.

Wickham, Hadley, Romain François, Lionel Henry, and Kirill Müller. 2020. *Dplyr: A Grammar of Data Manipulation.* https://CRAN.R-project.org/package=dplyr.

Xie, Yihui. 2020. *Knitr: A General-Purpose Package for Dynamic Report Generation in R.* https://CRAN.R-project.org/package=knitr.